

Processes

Operating Systems

```
top - 14:31:50 up 17 min, 1 user, load average: 0.38, 0.21, 0.17
Tasks: 243 total, 1 running, 242 sleeping, 0 stopped, 0 zombie
%Cpu0  :  0.0 us,  0.0 sy,  0.0 ni,100.0 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
%Cpu1  :  0.0 us,  0.0 sy,  0.0 ni,100.0 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
%Cpu2  :  0.0 us,  0.0 sy,  0.0 ni,100.0 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
%Cpu3  :  0.0 us,  0.0 sy,  0.0 ni,100.0 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
%Cpu4  :  0.0 us,  0.0 sy,  0.0 ni,100.0 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
%Cpu5  :  0.0 us,  0.0 sy,  0.0 ni,100.0 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
%Cpu6  :  0.0 us,  0.0 sy,  0.0 ni,100.0 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
%Cpu7  :  0.0 us,  0.0 sy,  0.0 ni,100.0 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
%Cpu8  :  0.0 us,  0.0 sy,  0.0 ni,100.0 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
%Cpu9  :  0.0 us,  0.0 sy,  0.0 ni,100.0 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
%Cpu10 :  0.0 us,  0.0 sy,  0.0 ni,100.0 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
%Cpu11 :  0.0 us,  0.0 sy,  0.0 ni,100.0 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
%Cpu12 :  0.0 us,  0.0 sy,  0.0 ni,100.0 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
%Cpu13 :  0.0 us,  0.0 sy,  0.0 ni,100.0 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
%Cpu14 :  0.0 us,  0.0 sy,  0.0 ni,100.0 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
%Cpu15 :  0.0 us,  0.0 sy,  0.0 ni,100.0 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
GiB Mem:  220.414 total,  0.846 used, 219.567 free,  0.015 buffers
GiB Swap:  0.000 total,  0.000 used,  0.000 free.  0.256 cached Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
1	root	20	0	34.2m	5.0m	3.6m	S	0.0	0.0	0:12.28	systemd
2	root	20	0	0.0m	0.0m	0.0m	S	0.0	0.0	0:00.00	kthreadd
3	root	20	0	0.0m	0.0m	0.0m	S	0.0	0.0	0:00.24	ksoftirqd/0
5	root	0	-20	0.0m	0.0m	0.0m	S	0.0	0.0	0:00.00	kworker/0:0H
6	root	20	0	0.0m	0.0m	0.0m	S	0.0	0.0	0:00.48	kworker/u128:0
7	root	20	0	0.0m	0.0m	0.0m	S	0.0	0.0	0:00.46	rcu_sched
8	root	20	0	0.0m	0.0m	0.0m	S	0.0	0.0	0:00.00	rcu_bh
9	root	20	0	0.0m	0.0m	0.0m	S	0.0	0.0	0:00.06	rcuos/0
10	root	20	0	0.0m	0.0m	0.0m	S	0.0	0.0	0:00.00	rcuob/0
11	root	rt	0	0.0m	0.0m	0.0m	S	0.0	0.0	0:00.42	migration/0
12	root	rt	0	0.0m	0.0m	0.0m	S	0.0	0.0	0:04.25	watchdog/0
13	root	rt	0	0.0m	0.0m	0.0m	S	0.0	0.0	0:06.80	watchdog/1
14	root	rt	0	0.0m	0.0m	0.0m	S	0.0	0.0	0:00.63	migration/1
15	root	20	0	0.0m	0.0m	0.0m	S	0.0	0.0	0:00.00	ksoftirqd/1
17	root	0	-20	0.0m	0.0m	0.0m	S	0.0	0.0	0:00.00	kworker/1:0H
18	root	20	0	0.0m	0.0m	0.0m	S	0.0	0.0	0:00.02	rcuos/1
19	root	20	0	0.0m	0.0m	0.0m	S	0.0	0.0	0:00.00	rcuob/1
20	root	rt	0	0.0m	0.0m	0.0m	S	0.0	0.0	0:00.00	watchdog/2
21	root	rt	0	0.0m	0.0m	0.0m	S	0.0	0.0	0:00.00	migration/2
22	root	20	0	0.0m	0.0m	0.0m	S	0.0	0.0	0:00.00	ksoftirqd/2
24	root	0	-20	0.0m	0.0m	0.0m	S	0.0	0.0	0:00.00	kworker/2:0H
25	root	20	0	0.0m	0.0m	0.0m	S	0.0	0.0	0:00.00	rcuos/2
26	root	20	0	0.0m	0.0m	0.0m	S	0.0	0.0	0:00.00	rcuob/2
27	root	rt	0	0.0m	0.0m	0.0m	S	0.0	0.0	0:00.10	watchdog/3
28	root	rt	0	0.0m	0.0m	0.0m	S	0.0	0.0	0:00.00	migration/3

executable format

defines the structure of
an executable stored on
a disk or in memory

usually produced by a
program linker

executable image

executable format header

program header

<code>

<global data>

<static data>



The diagram illustrates the structure of an executable image. It is represented as a large rectangle with a light blue border. Inside, the components are listed vertically: 'executable format header', 'program header', '<code>', '<global data>', and '<static data>'. The 'program header' is connected to the '<code>', '<global data>', and '<static data>' sections by three curved arrows pointing to the left. The '<static data>' section is also connected to the 'program header' by a curved arrow pointing to the right.

common executable formats

PE (Windows)

ELF (Linux...)

Mach-O (Mac OS X)

process descriptor

process structure

process identification

process state

process priority

...

cpu state

virtual memory

opened files

...

code

global data

static data

...

parent

working directory

exit status

...

cpu state

registers

program counter

stack pointer

...

memory map

process page table

kernel maintains a list
of process descriptors.

process structure


process identification
process state
process priority
...
cpu state
virtual memory
opened files
...
code
global data
static data
...
parent
working directory
exit status
...

process structure

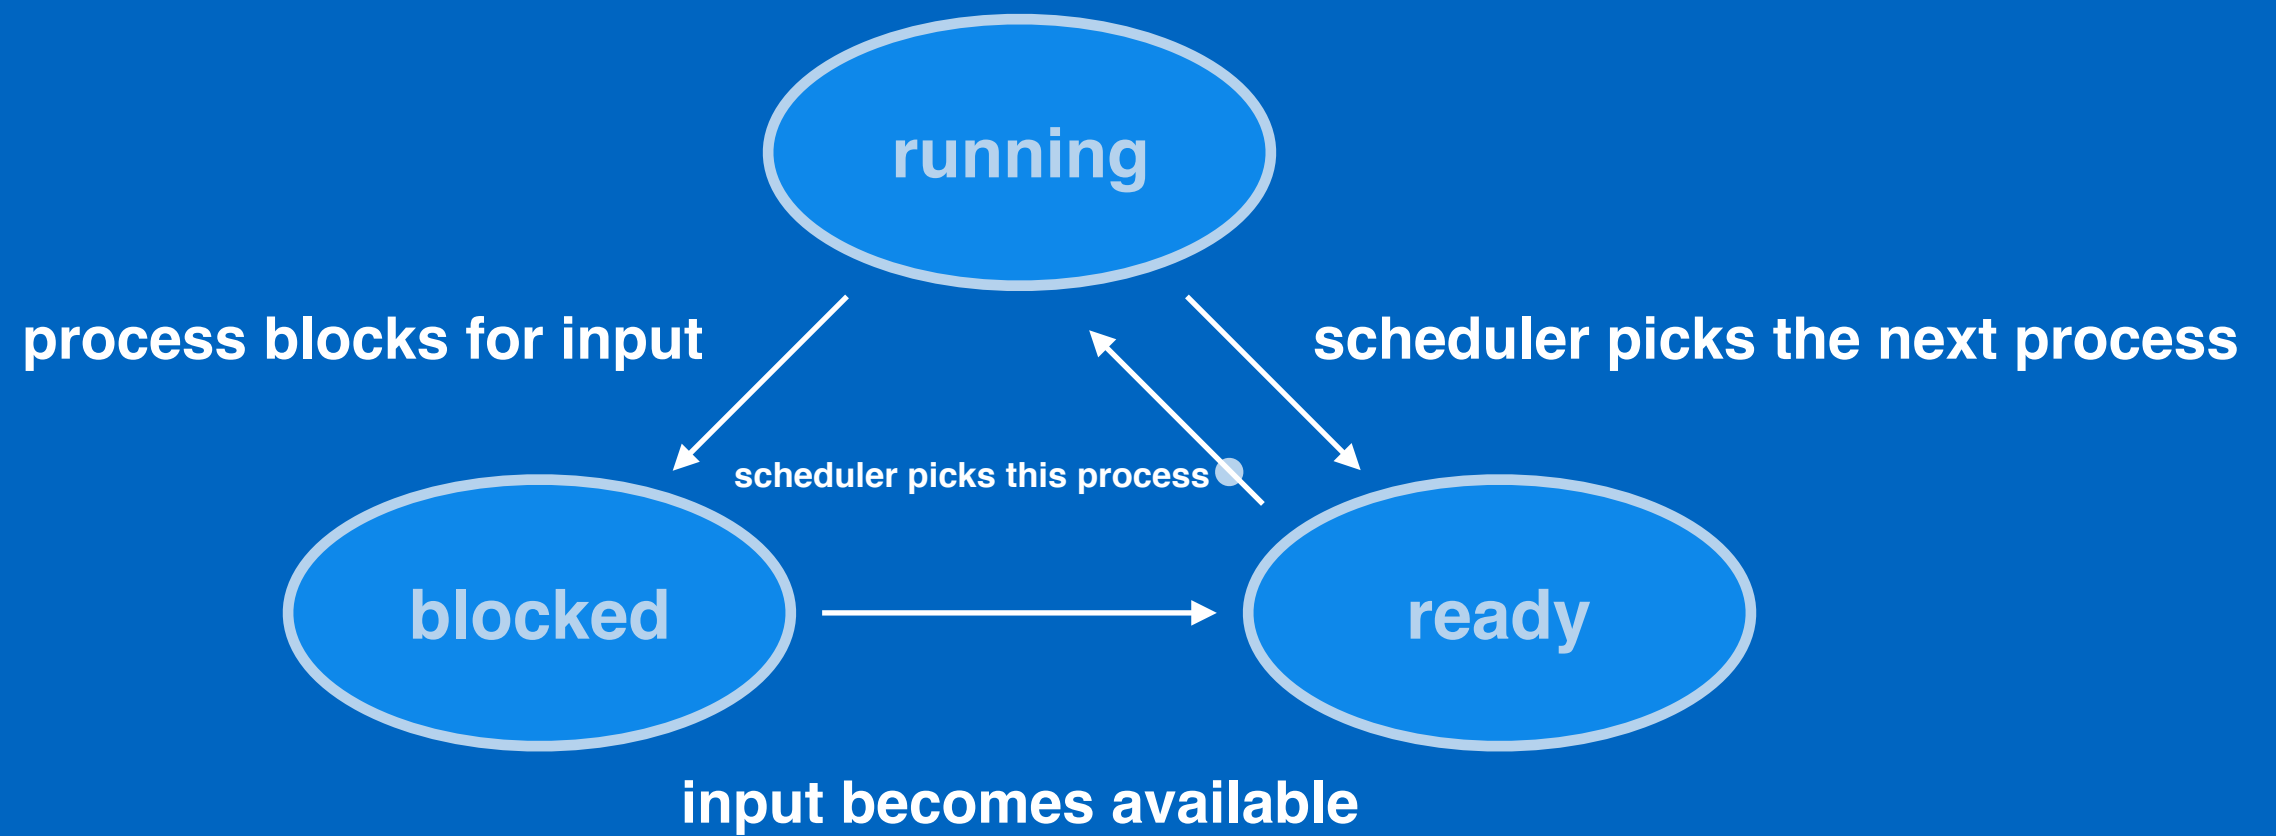
process identification
process state
process priority
...
cpu state
virtual memory
opened files
...
code
global data
static data
...
parent
working directory
exit status
...

process structure

process identification
process state
process priority
...
cpu state
virtual memory
opened files
...
code
global data
static data
...
parent
working directory
exit status
...



process states



context switch

reasons to switch
between executable units

a new process is created
the process exits

the process blocks
waits for input/output
waits to enter a critical section

a hardware interrupt occurs
input/output device interrupts
clock interrupts

steps to switch a context

1. save the current CPU state to memory

process structure

process identification
process state
process priority

cpu state
virtual memory
opened files

code
global data
static data

parent
working directory
exit status

process structure

process identification
process state
process priority

...
cpu state
virtual memory
opened files

...
code
global data
static data

...
parent
working directory
exit status

...


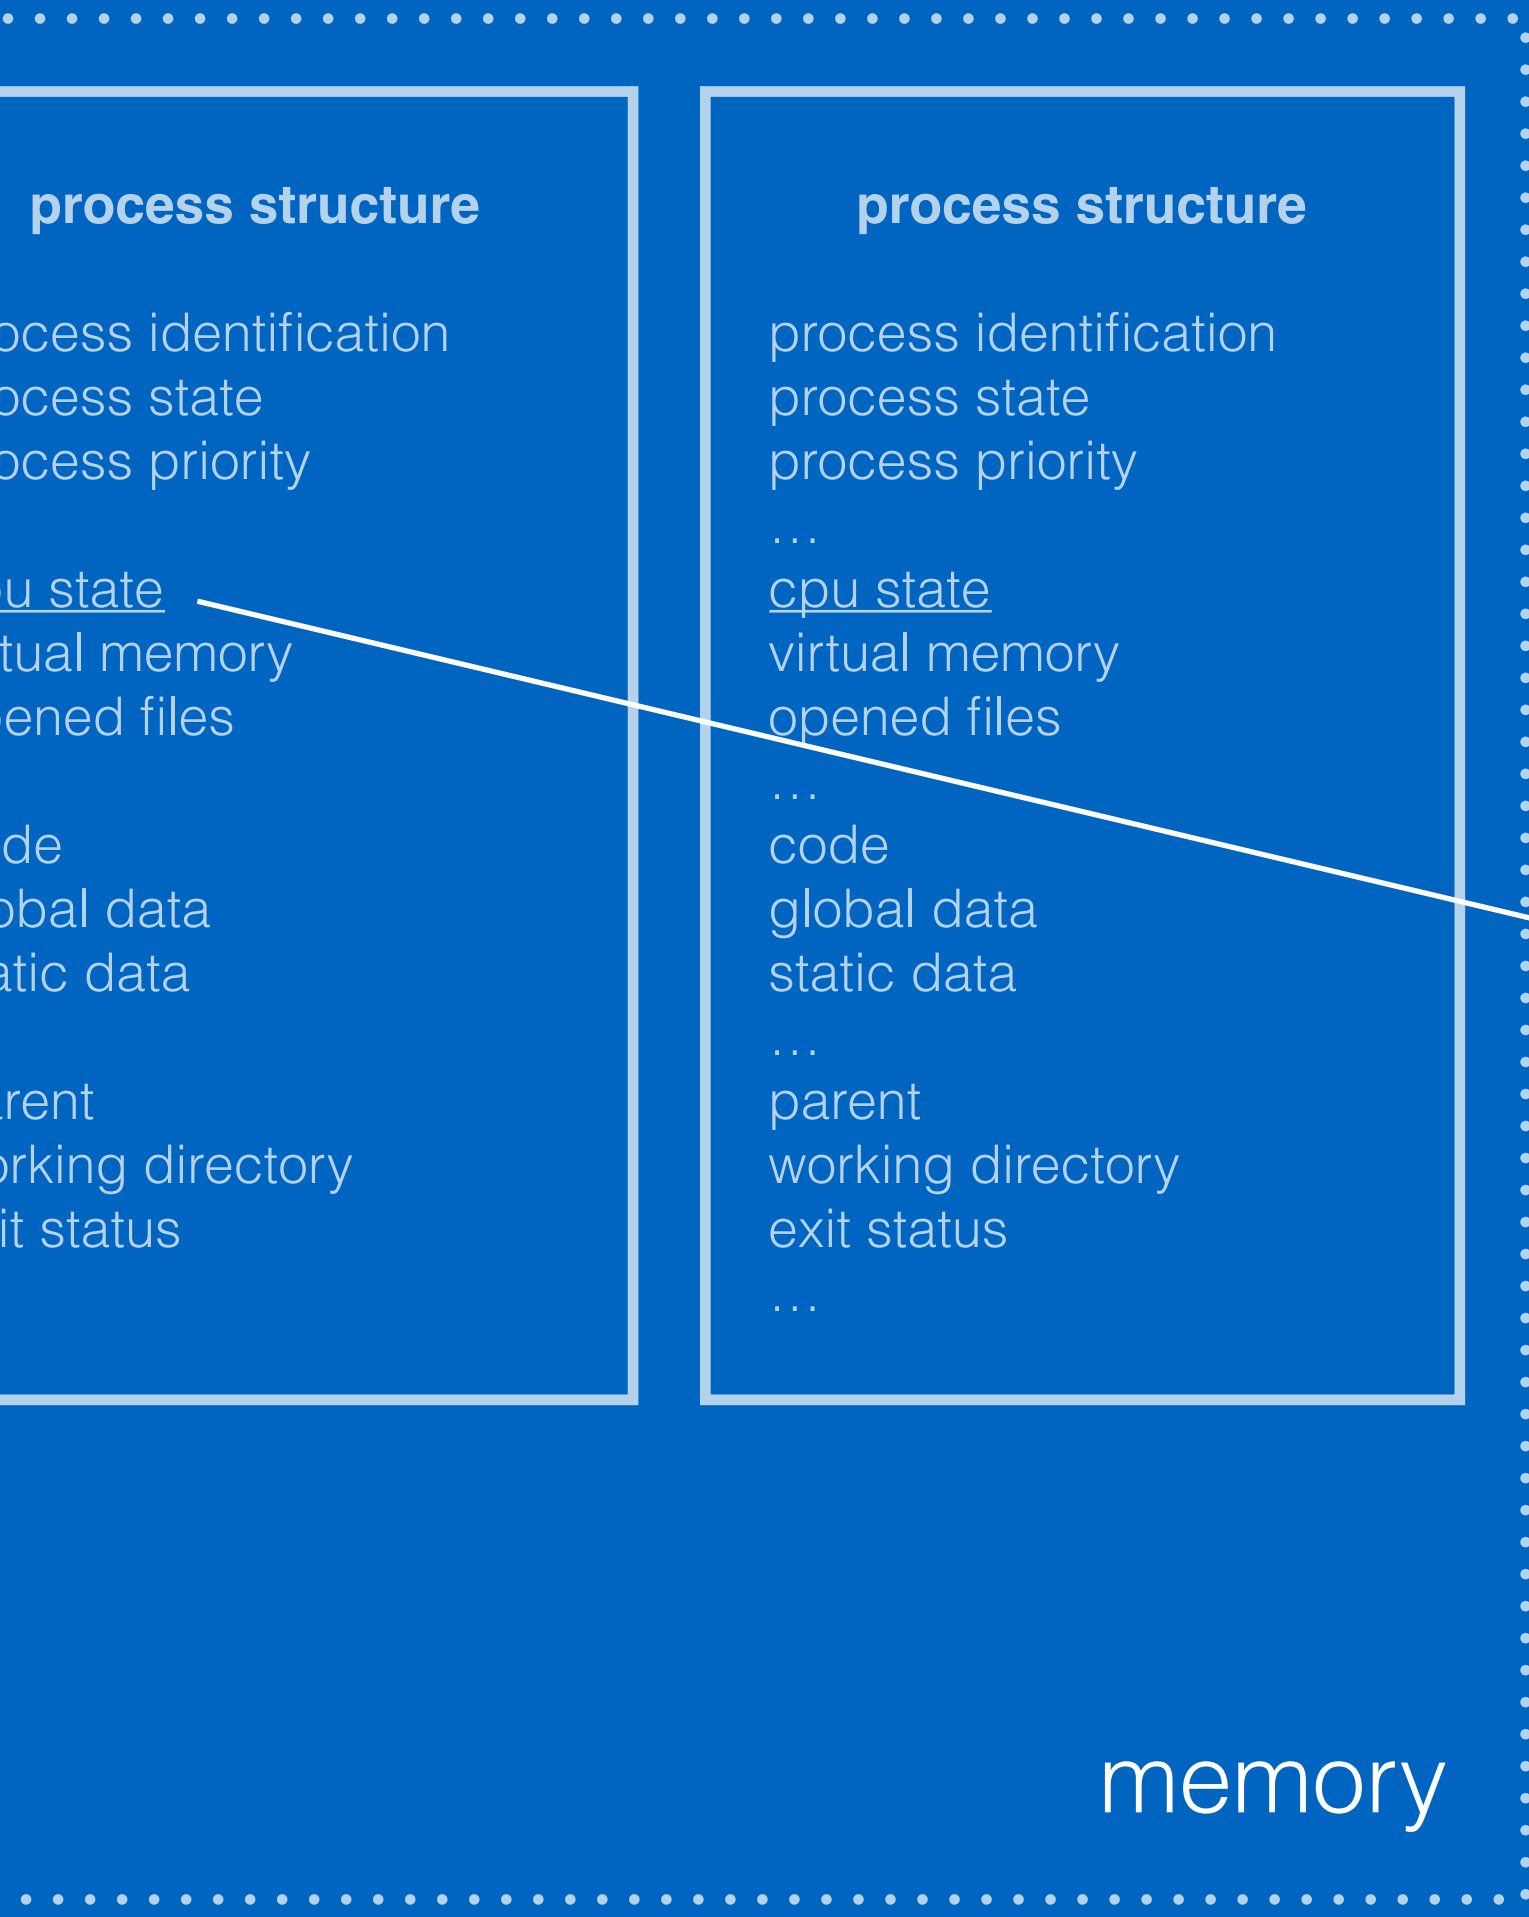
CPU

reg0
reg1
instruction pointer
stack pointer

...

memory

2. load a CPU state for the next executable unit from memory



threads

thread (or a lightweight process) is an executable unit that shares certain elements with its owning process.

shared memory address space

shared opened files

shared global data

...

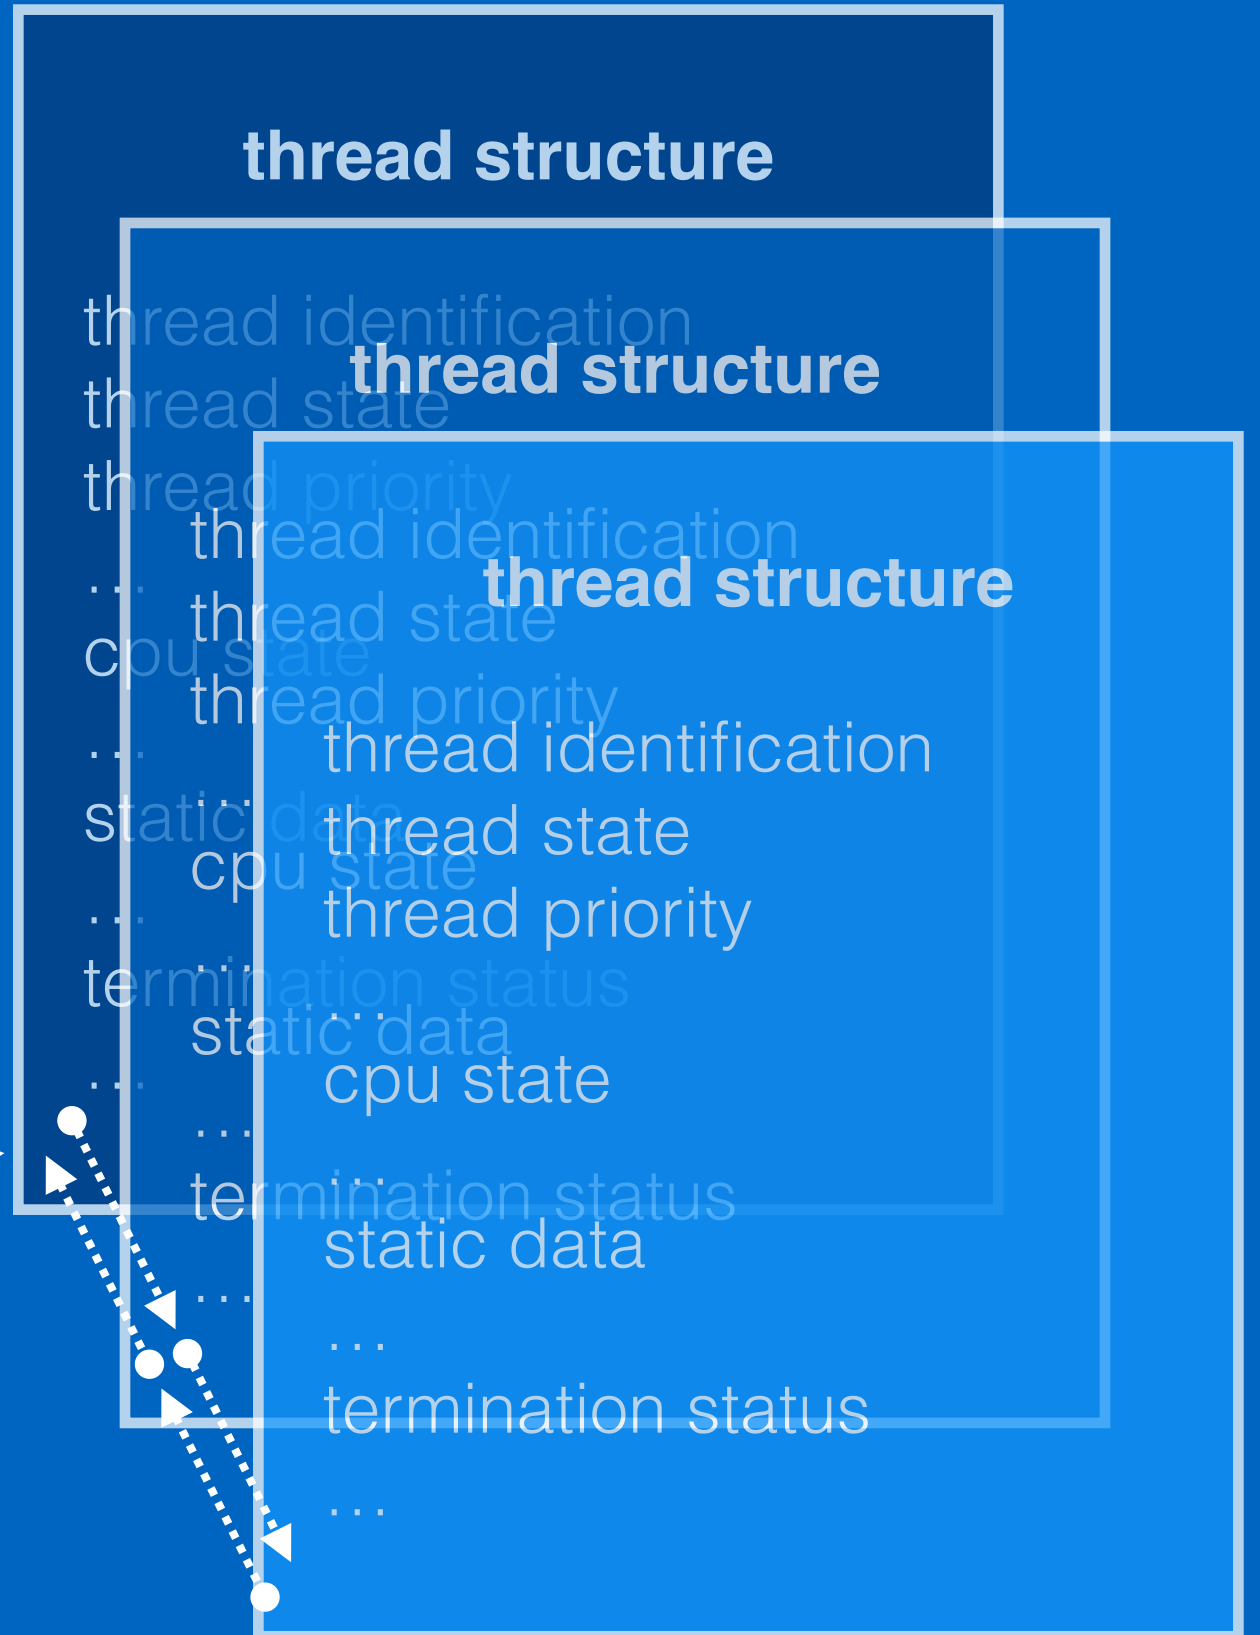
private state
private registers
private stack
...

process structure

process identification
process state
process priority
...
cpu state
virtual memory
opened files
...
code
global data
static data
...
parent
threads •
working directory
exit status
...

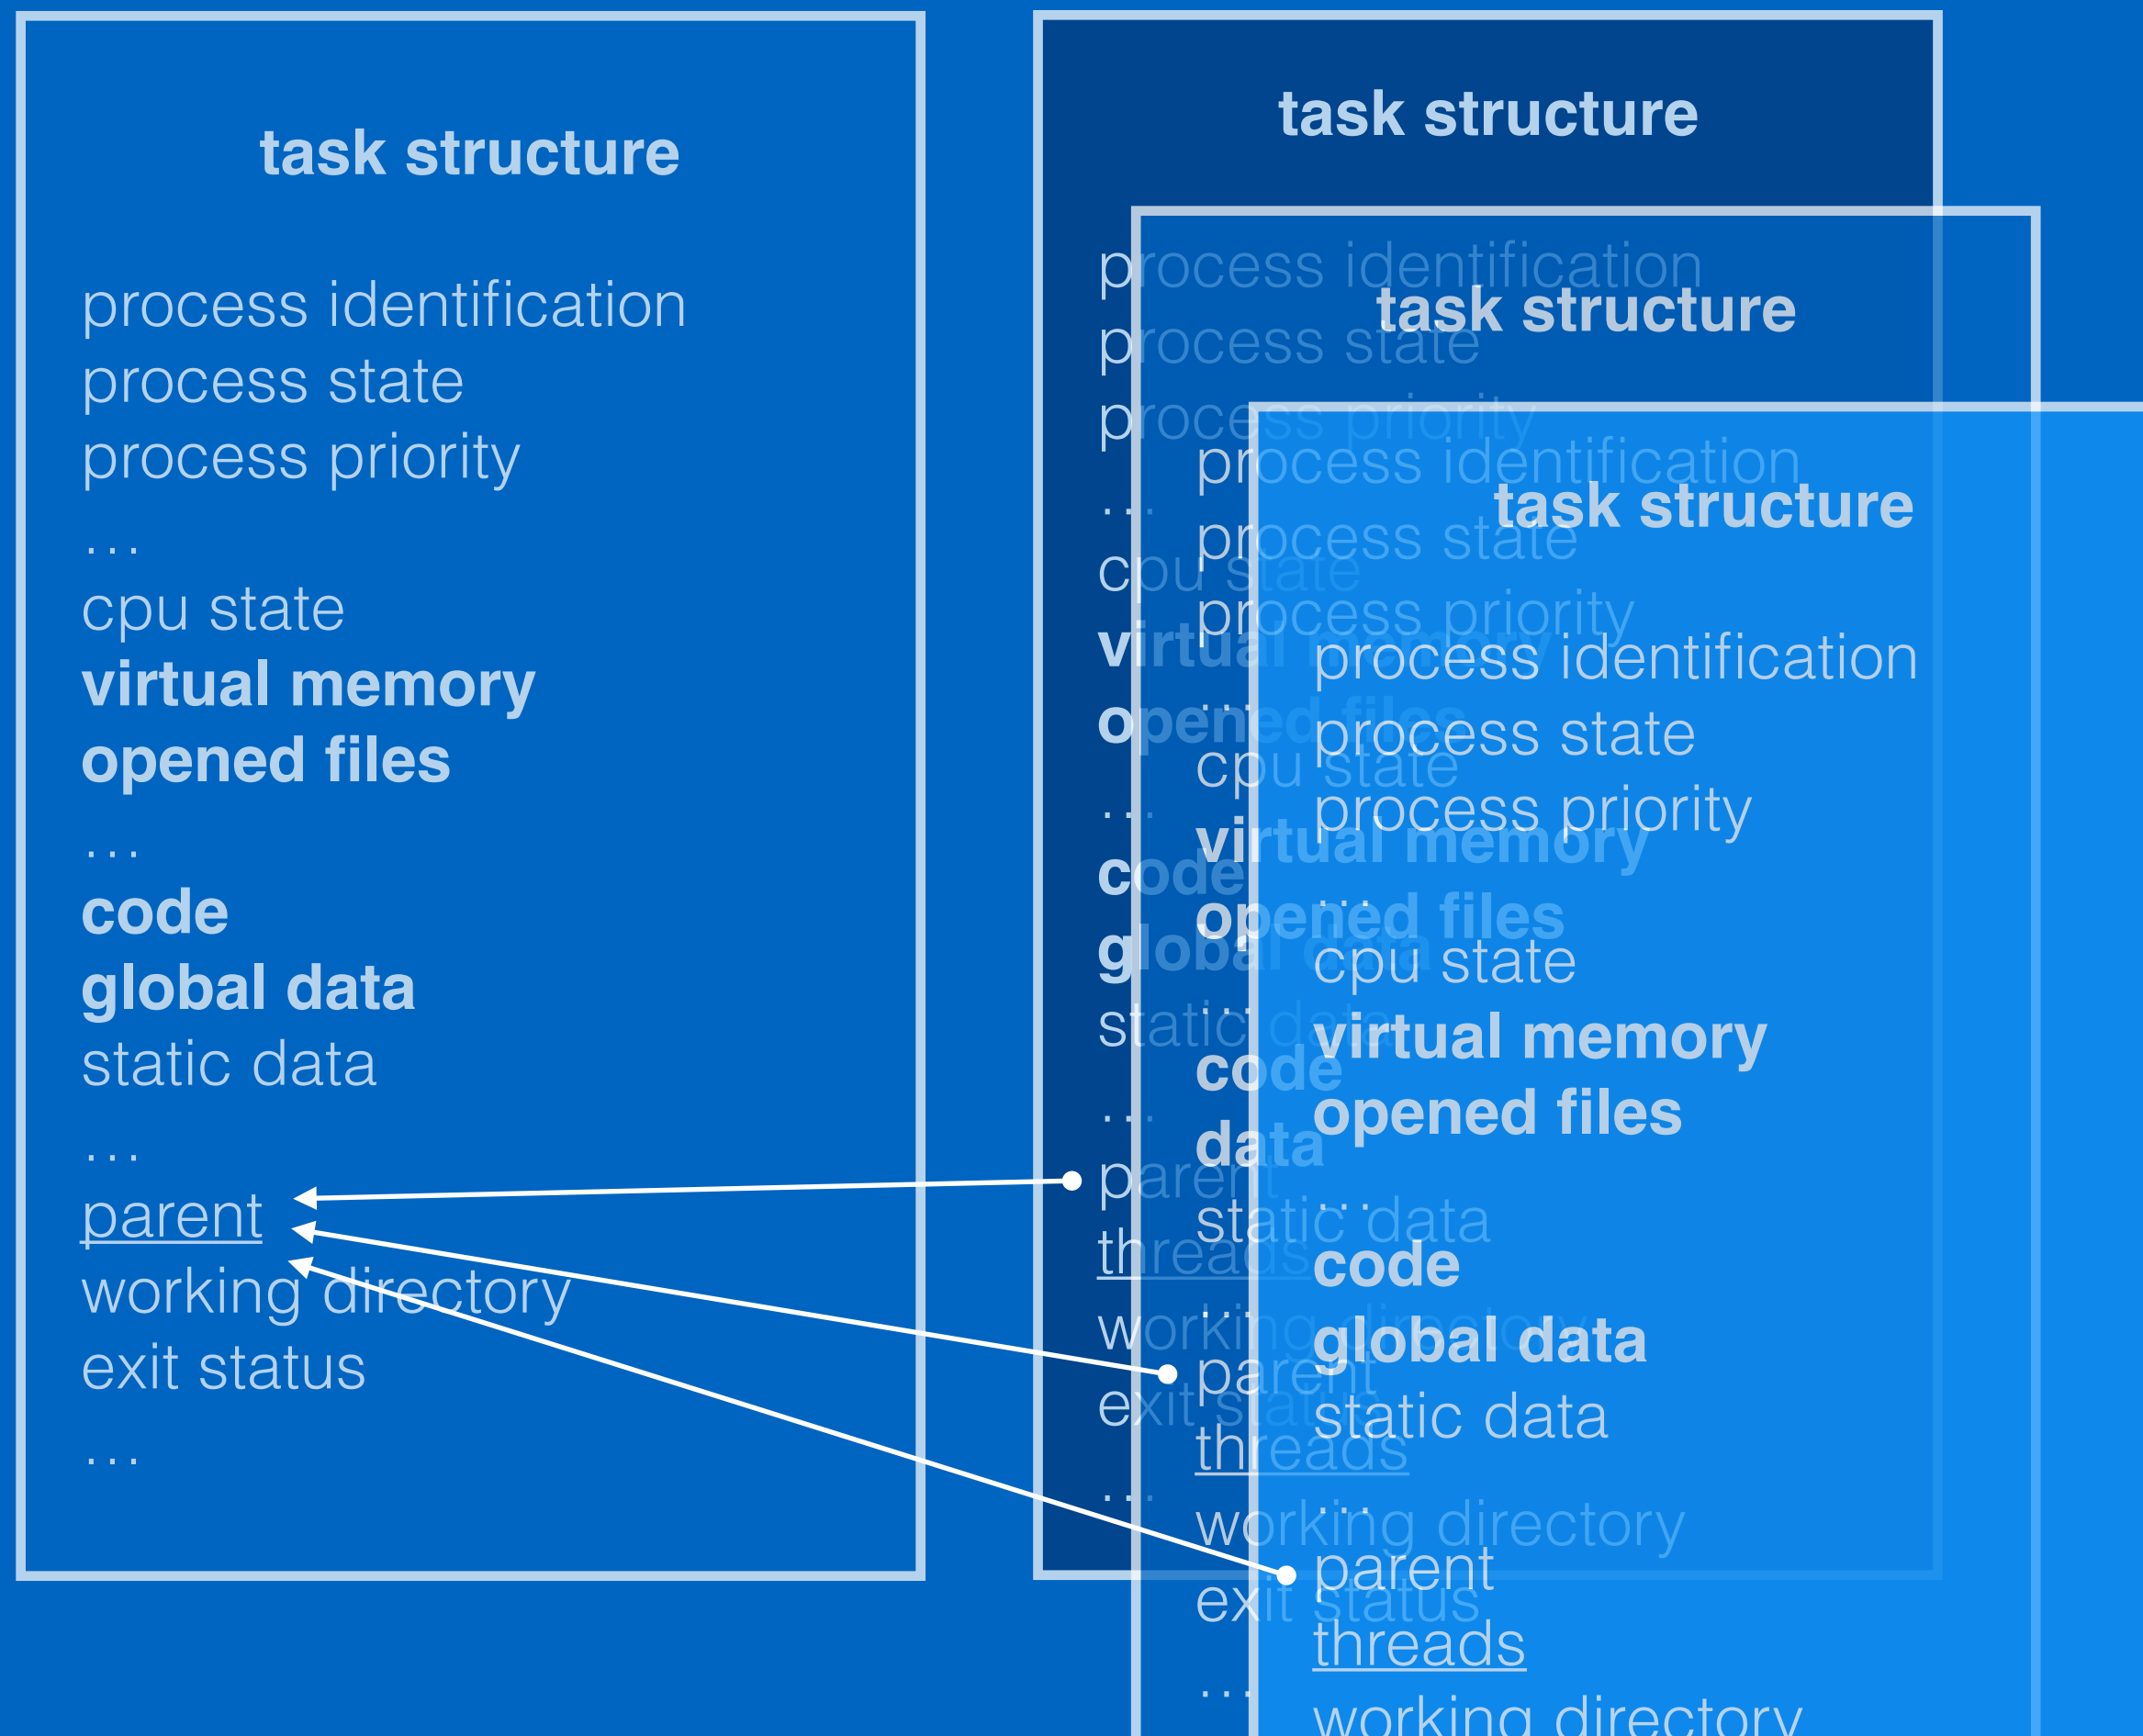
thread structure

thread identification
thread state
thread priority
thread identification
thread state
thread priority
thread identification
thread state
thread priority
static data
cpu state
termination status
static data
cpu state
termination status
static data
...
termination status
...



same kernel data structure
can be used for process
descriptors and threads.

such approach simplifies
scheduling and context
switching logic.



Reading

Operating Systems Design and Implementation,
Third Edition by Andrew S. Tanenbaum

Chapter 2 (2.1, 2.4)

Supplemental Reading

Understanding the Linux kernel, Third Edition by
Daniel P. Bovet and Marco Cesati

Chapter 3

Supplemental Reading

Windows Internals, Part 1 (6th Edition) by Mark E. Russinovich and David A. Solomon

Chapter 5

Supplemental Reading

Mac OS X and iOS internals : to the apple's core
by Jonathan Levin

Chapter 4