

Contents

Midterm Notes	1
Labs	1
true	1
false	1
yes	1
echo	2
cat	2
pwd	2
touch	2
mkdir	2
mv	2
rm	2
cp	3
ls	3
Kernel Project	3
get_pids.c	3
get_task_info.c	3
task_info.h	3
Presentations	3
Introduction	3
History	4
Boot Process	5
Processes	5
Scheduling	6

Midterm Notes

Labs

- to see the return code: `echo $?`
- for syscall numbers `/usr/include/asm/unistd_64.h`

true

- returns 0 (C false)

false

- returns 1 (C true)

yes

- prints y or provided arguments separated by spaces forever
- use `print`, `strlen`

```
long strlen(const char *s) {  
    long i;
```

```

    for (i = 0; s[i] != '\0'; i++) ;
    return i;
}

```

echo

- print space-separated arguments, ending in a newline

cat

- `openat` for file
- for each argument:

```

ssize_t bytes_read;
while((bytes_read = read(fd, buff, BUFF_SIZE)) > 0) {
    write(1, buff, bytes_read);
}
close(fd);

```

pwd

- `getcwd` syscall

touch

- if file does not exist: create empty file
 - `openat` with arguments – `O_WRONLY|O_CREAT|O_NOCTTY|O_NONBLOCK`
- if file exists – update time stamp
 - `openat` with arguments – `O_WRONLY|O_CREAT|O_NOCTTY|O_NONBLOCK`

mkdir

- `mkdir` with 0777 as argument
- user, group, other

mv

- `renameat2(AT_FDCWD, "test/", AT_FDCWD, "test-2/", RENAME_NOREPLACE)`
= 0
- this is done for any type
- move all but the last file into the last location

rm

- for `rm -rf`

```

unlinkat(AT_FDCWD, "test-2/", AT_REMOVEDIR) = 0

```

- for `rm`

```
unlinkat(AT_FDCWD, "touch_2.c", 0) = 0
```

cp

- if `cp -r`
 - check if dir exists with `fsstatat`
 - * if not, `mkdir`
 - for each file in the directory – open the file, read from it, then write to it
 - all with `openat`
- if just `cp`
 - stat both files to check existence
 - read from one file, write to the other one

ls

- `openat(AT_FDCWD, ".", O_RDONLY|O_NONBLOCK|O_CLOEXEC|O_DIRECTORY) = 3`
- then `fstat(3, {st_mode=S_IFDIR|0755, st_size=4096, ...}) = 0`
- `getdents64(3, 0x55dc15be8420 /* 9 entries */, 32768) = 224`
- then simply write the `names[]` to `stdout`

Kernel Project

get_pids.c

- see file

get_task_info.c

- see file

task_info.h

- see file

Presentations

Introduction

- most fundamental system program is the operating system
- it manages the computer's resources and provides a base for application development
- Hardware → Kernel → Applications → User
- it does resource management and machine abstraction
- resources:
 - CPU
 - memory
 - disks
 - other devices
- machine abstraction:
 - processes
 - threads

- files
- shell
- process, files systems, device abstraction can be modeled as trees
- system calls hide the complexity of the real machine
- single-tasking, multi-tasking, single-user, multi-user are what they sounds like
- real-time operating system – should not have a time delay on the input
- distributed – separate nodes that are networked
- embedded os – specific tasks for non-pcs
- library os – single address-space machine that has the bare minimum of stuff for programs to run and they are then compiled and run directly on hardware without an os
- microkernel: bare minimum runs in kernel space, the rest in user mode
- hybrid kernel: mix of things – Windows, OS X
- monolithic kernel: everything but applications run in kernel mode – Linux

History

1945 to 1955

- plug boards, punch cards

1955 to 1965

- transistors
- assembly languages
- programming languages
- jobs, job control systems, batch systems (jobs without user input than can be run as resources permit)

1965 to 1980

- integrated circuits
- major os projects: IBM OS/360, CTSS, MULTICS
- multi threading
- time sharing
- UNIX 1970, Ken Thompson, Dennis Ritchie
- CTSS -> Multics -> UNIX
- first in asm, then later in C

1980 to 1990

- large scale integration circuits – CPUs
- microcomputers
- PCs
- DOS (apple, microsoft,...)
- GNU, BSD

1990 to present

- internet

- mobile boom
- iot

Boot Process

1. power up
2. internal firmware
 - BIOS
 - UEFI
 - custom stuff
3. boot loader
 - multiple stages
 - MBR, GPT
4. kernel
 - CPU and hardware setup
5. initialization service
 - init, upstart, systemd, launchd
6. user space programs
 - login, shell

Processses

- executable format: structure of an exe
- produced by a linker
- format header, program header, then links to code, global and static data
- common formats: PE, ELF, Mach-O
- each process has: pid, state, priority, cpu state (registers, program counter, stack pointer), virtual memory (memory map), opened files, code, data, parent id, working directory, exit status,...
- the kernel maintains a list of all the process descriptors
- process states: running, ready, blocked for scheduling
- context switching to give all exes an equal amount of time
- processes can be created and exit when they are done
- io devices can interrupt, the clock can interrupt

Context Switching

1. save current cpu state
2. load a CPU state from next executable
3. run that one

Threads

- lightweight process that shares some elements with its parent process
- e.g. memory address space, opened files, global data, ...
- they tend to have a private stack, registers, state
- same data structure for processes and threads

Scheduling

- either compute-bound or IO-bound processes
- the time-bottleneck is either computation or IO
- compute-bound: long CPU bursts, little IO
- IO-bound: short CPU bursts, a lot of IO
- a scheduler can start and stop these processes to make a efficient use of the system resources
- preemptive scheduling – mostly by priority
- non-preemptive goes by what processes finish when
- fairness and balance are common goals

Batch Scheduler

- maximize throughput
- minimize turnabout
- maximize CPU utilization
- first-come first-served
- shortest job first
- shortest remaining time next

Interactive Scheduler

- minimize response time
- round-robin scheduling
- shortest process next
- lottery scheduling
- fair-share scheduling

Real-time Scheduler

- meet deadlines
- hard real time vs soft real time:
 - hard real time have deadlines that must be met
 - soft real time has a certain tolerance but still not desirable
- periodic or non-periodic events